



# Kubernetes Best Practices

Sandeep Dinesh  
Developer Advocate  
[@sandeepdinesh](#)  
[github.com/thesandlord](#)



# Kubernetes is really flexible

But you might 🔫 yourself in the 👟



# Building Containers

Don't trust arbitrary base images!

# Static Analysis of Containers



<https://github.com/coreos/clair>

<https://github.com/banyanops/collector>

# Use small base images

# Overhead

## Node.js App

Your App → 5MB

Your App's Dependencies → 95MB

Total App Size → 100MB

## Docker Base Images:

node:8 → 667MB

node:8-wheezy → 521MB

node:8-slim → 225MB

node:8-alpine → 63.7MB

scratch → ~50MB



# Overhead

## Node.js App

Your App → 5MB

Your App's Dependencies → 95MB

Total App Size → 100MB

← 6.6x App Size!!

## Docker Base Images:

node:8 → 667MB

node:8-wheezy → 521MB

node:8-slim → 225MB

node:8-alpine → 63.7MB

scratch → ~50MB

# Overhead

## Node.js App

Your App → 5MB

Your App's Dependencies → 95MB

Total App Size → 100MB

## Docker Base Images:

node:8 → 667MB

node:8-whispery → 521MB

node:8-slim → 225MB

node:8-alpine → 63.7MB

scratch → ~50MB

← 6.6x App Size!!

← 13.3x "min" overhead!!

# Overhead

## Node.js App

Your App → 5MB

Your App's Dependencies → 95MB

Total App Size → 100MB

## Docker Base Images:

node:18 → 667MB

node:8-whispery → 521MB

node:8-slim → 225MB

node:8-alpine → 63.7MB

scratch → ~50MB

← 6.6x App Size!!

← 13.3x "min" overhead!!

## Pros:

Builds are faster

Need less storage

Cold starts (image pull) are faster

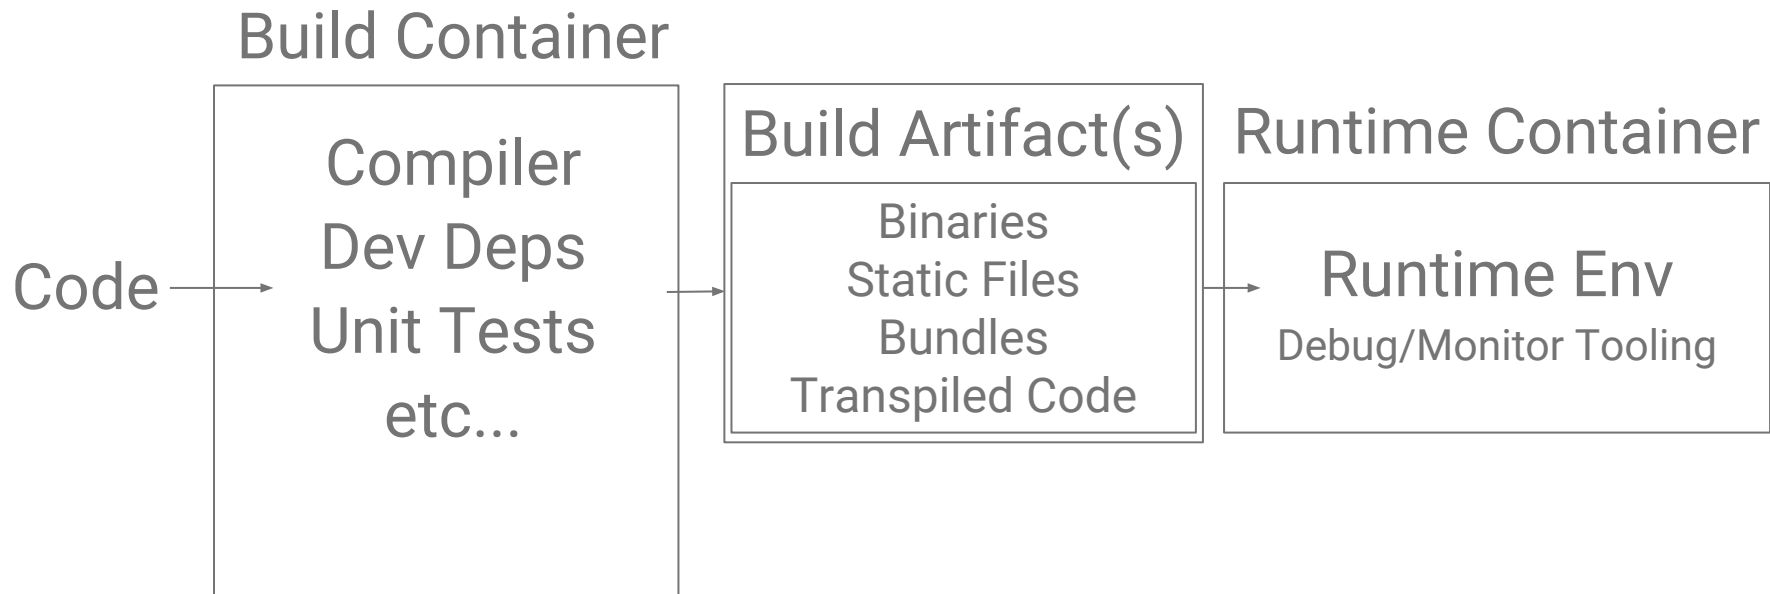
Potentially less attack surface

## Cons:

Less tooling inside container

"Non-standard" environment

Use the “builder pattern”



# Docker bringing native support for multi-stage builds in Docker CE 17.05



# Container Internals

Use a non-root user inside the container



# Example Dockerfile

```
FROM node:alpine  
RUN apk update && apk add imagemagick
```

```
RUN groupadd -r nodejs
```

```
RUN useradd -m -r -g nodejs nodejs
```

```
USER nodejs
```

```
ADD package.json package.json
```

```
RUN npm install
```

```
ADD index.js index.js
```

```
CMD npm start
```

# Enforce it!

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world
spec:
  containers:
    # specification of the pod's containers
    # ...
  securityContext:
    runAsNonRoot: true
```

# Make the filesystem read-only

# Enforce it!

```
apiVersion: v1
kind: Pod
metadata:
  name: hello-world
spec:
  containers:
    # specification of the pod's containers
    # ...
  securityContext:
    runAsNonRoot: true
    readOnlyRootFilesystem: true
```

# One process per container

Don't restart on failure. Crash cleanly instead.

# Log to stdout and stderr

Add “dumb-init” to prevent zombie processes



# Example Dockerfile

```
FROM node:alpine
RUN apk update && apk add imagemagick

RUN groupadd -r nodejs
RUN useradd -m -r -g nodejs nodejs
USER nodejs

ADD https://github.com/Yelp/dumb-init/releases/download/v1.2.0/dumb-init_1.2.0_amd64 \
    /usr/local/bin/dumb-init
RUN chmod +x /usr/local/bin/dumb-init
ENTRYPOINT ["/usr/bin/dumb-init", "--"]

ADD package.json package.json
RUN npm install
ADD index.js index.js
CMD npm start
```

Good News: No need to do this in K8s 1.7



# Deployments

Use the “record” option for easier rollbacks

```
$ kubectl apply -f deployment.yaml --record
```

```
...
```

```
$ kubectl rollout history deployments my-deployment
```

```
deployments "ghost-recorded"
```

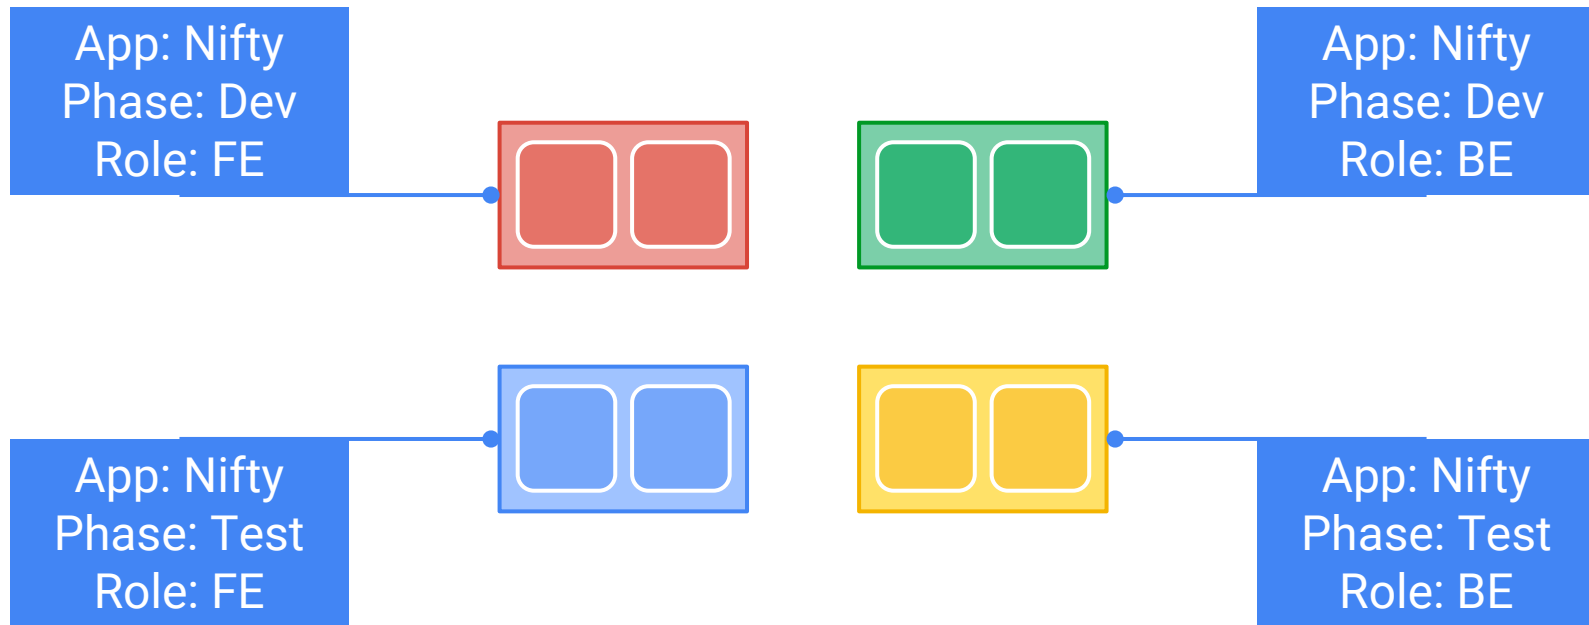
REVISION	CHANGE-CAUSE
1	kubectl apply -f deployment.yaml --record
2	kubectl edit deployments my-deployment
3	kubectl set image deployment/my-deplyoment my-container=app:2.0

# Use plenty of descriptive labels

# Labels

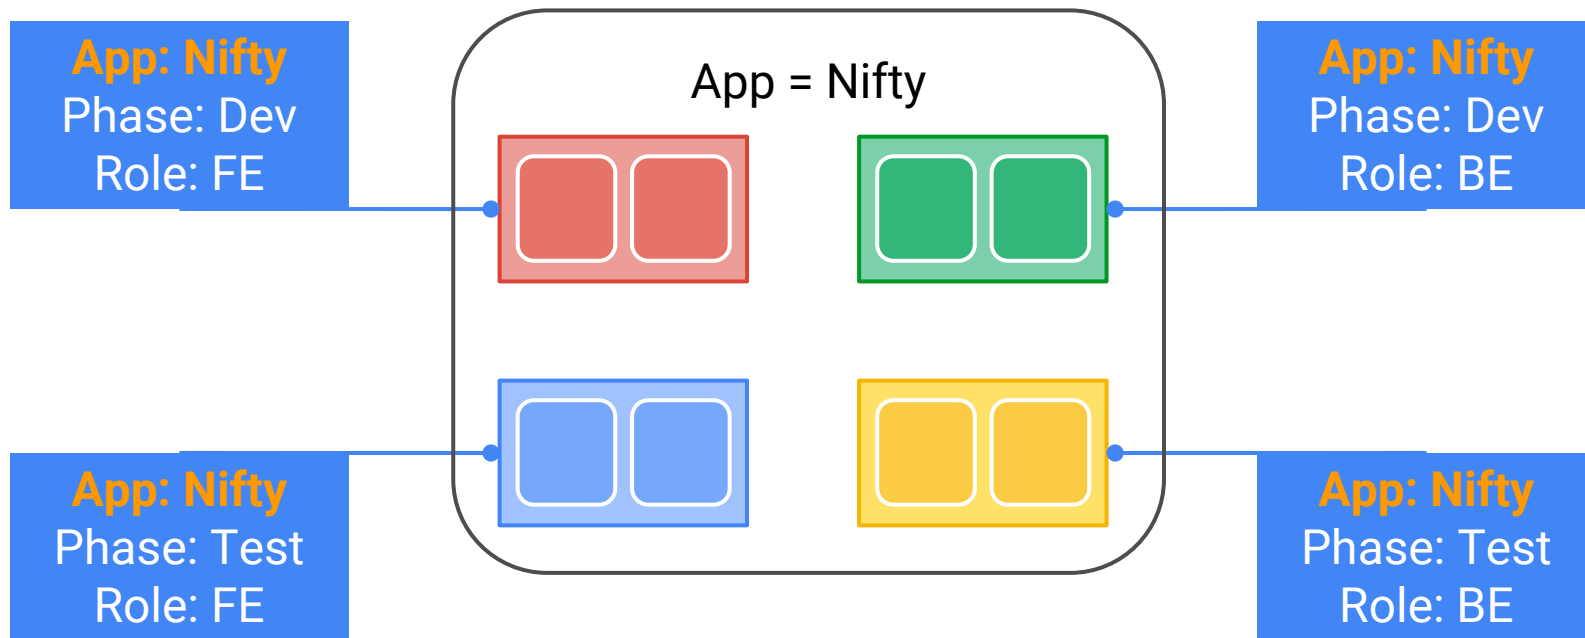
```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: web
spec:
  replicas: 12
  template:
    metadata:
      labels:
        name: web
        color: blue
        experimental: 'true'
```

# Labels

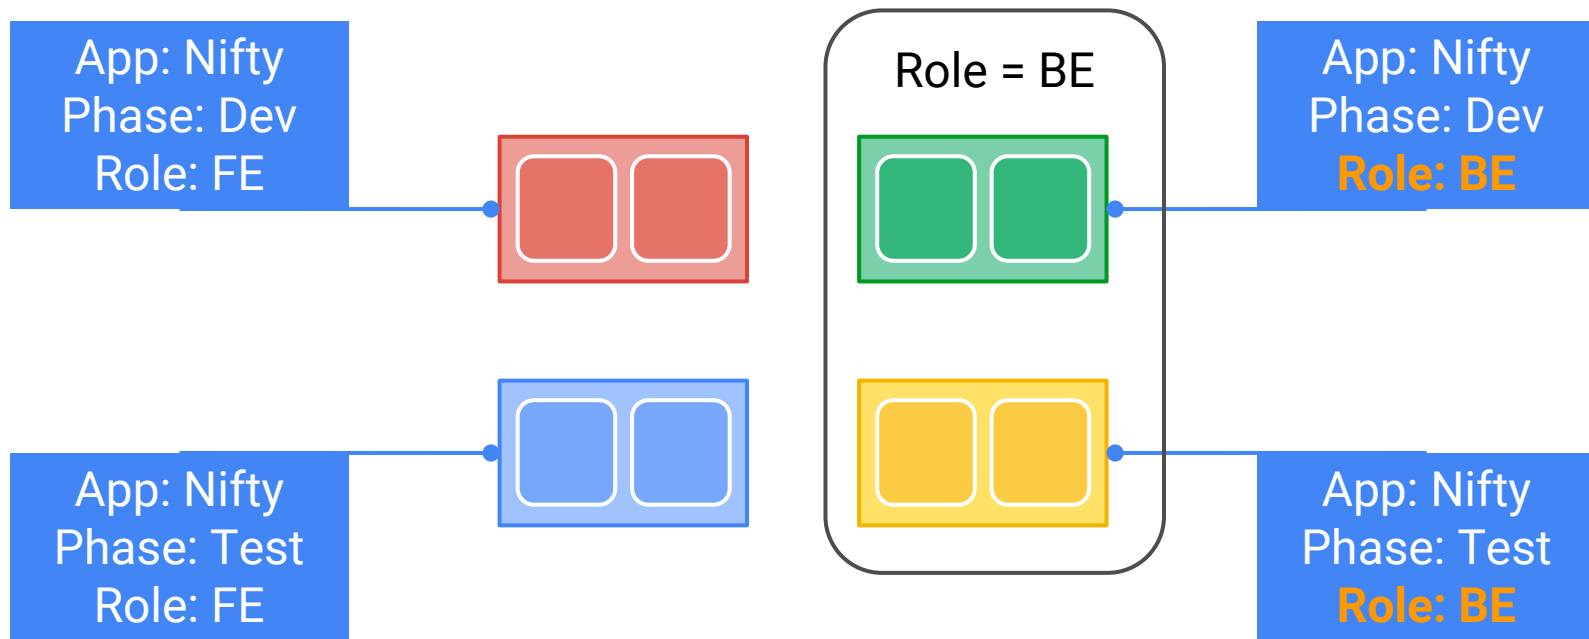




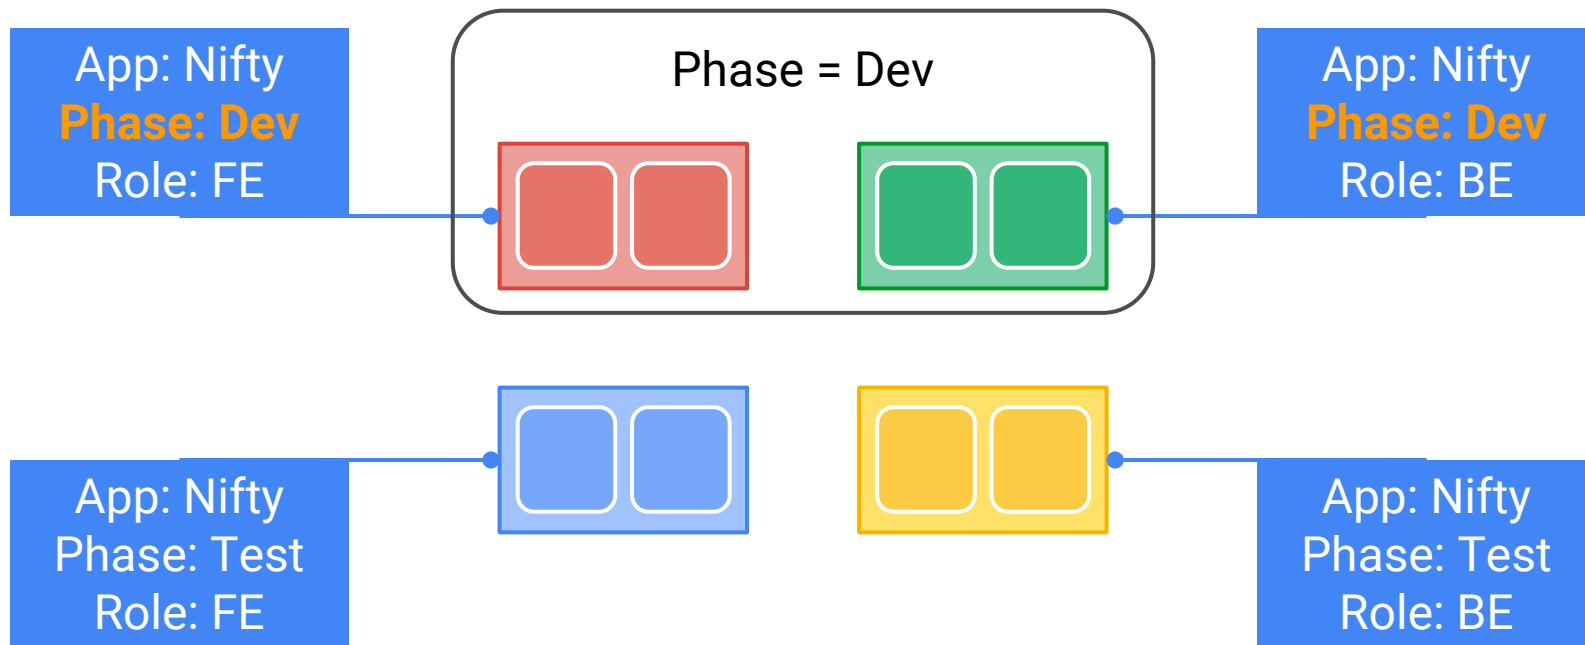
# Labels



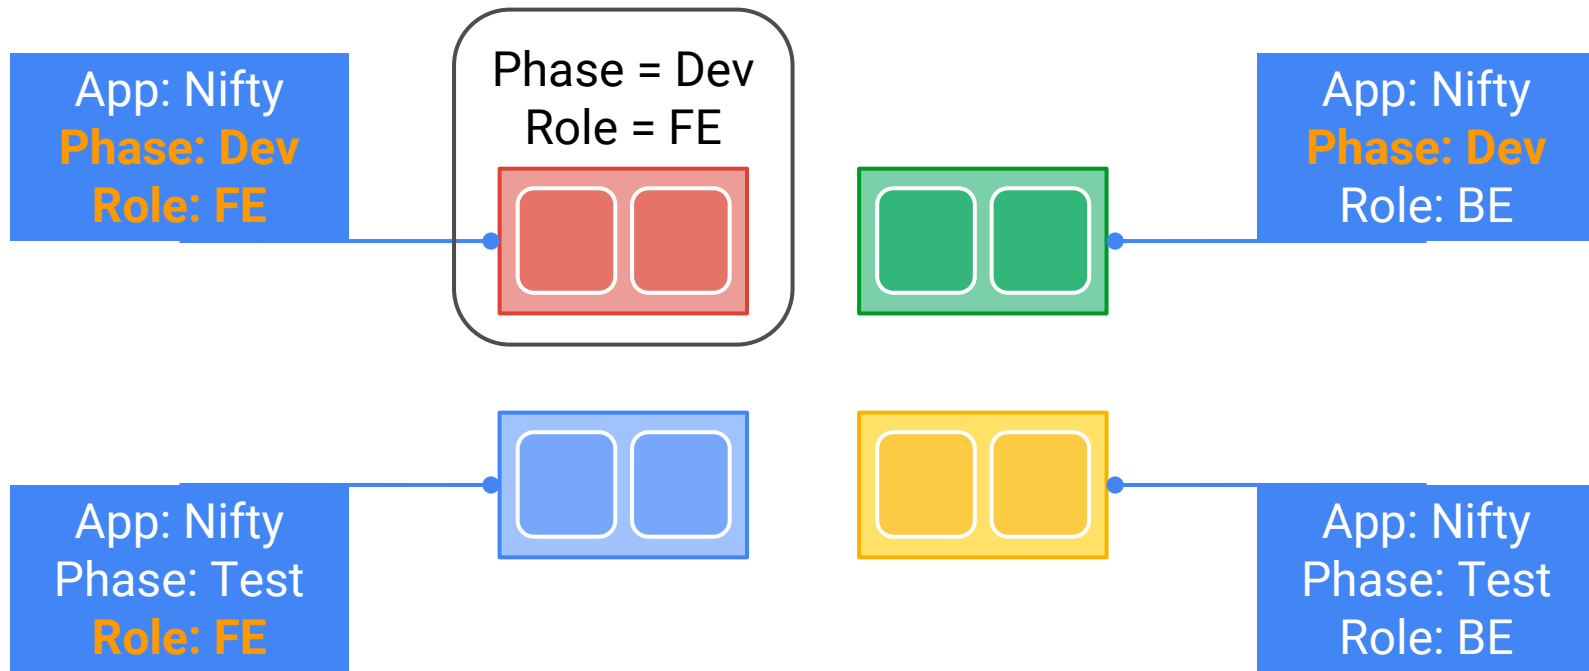
# Labels



# Labels

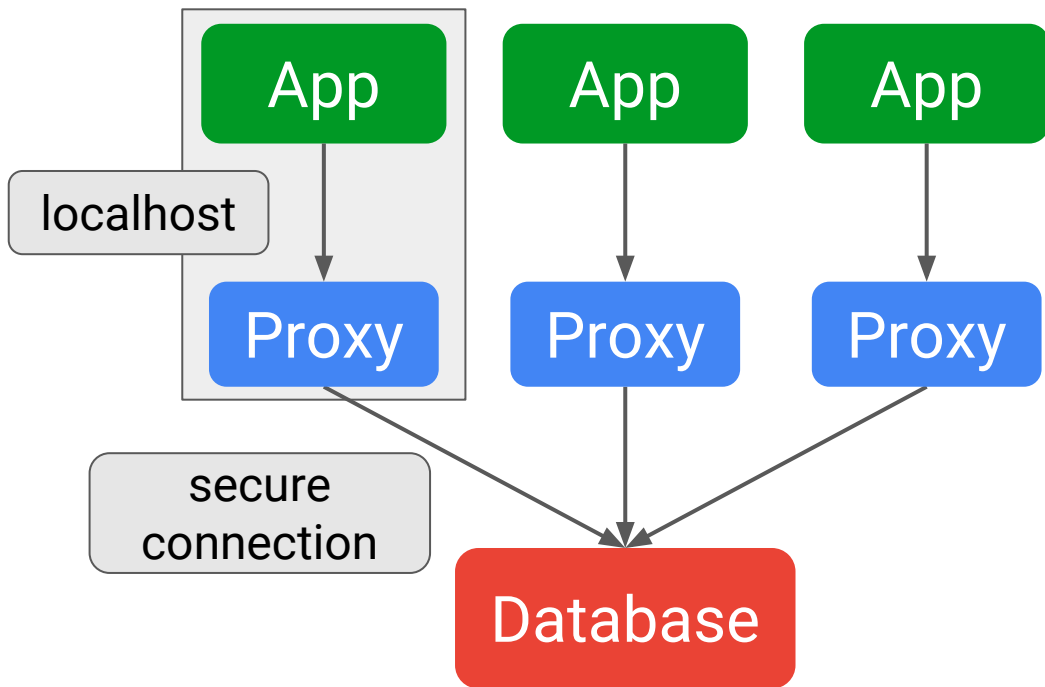


# Labels

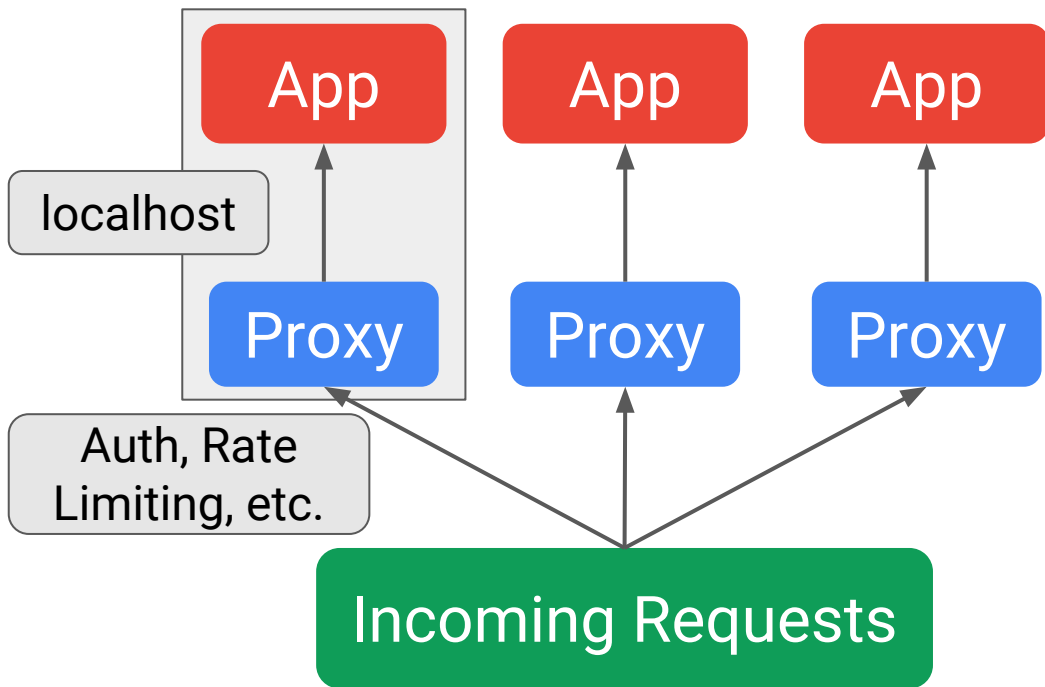


Use sidecar containers for proxies, watchers, etc

# Examples



# Examples



Don't use sidecars for bootstrapping!



Use init containers instead!

```
apiVersion: v1
kind: Pod
metadata:
  name: awesomeapp-pod
  labels:
    app: awesomeapp
  annotations:
    pod.beta.kubernetes.io/init-containers '[
      {
        "name": "init-myapp",
        "image": "busybox",
        "command": ["sh", "-c", "until nslookup myapp; do echo waiting for myapp; sleep 2; done;"]
      },
      {
        "name": "init-mydb",
        "image": "busybox",
        "command": ["sh", "-c", "until nslookup mydb; do echo waiting for mydb; sleep 2; done;"]
      }
    ]'
spec:
  containers:
  - name: awesomeapp-container
    image: busybox
    command: ['sh', '-c', 'echo The app is running! && sleep 3600]
```

Don't use `:latest` or no tag

Readiness and Liveness probes are your friend

# Health Checks

Readiness → Is the app ready to start serving traffic?

- Won't be added to a service endpoint until it passes
- Required for a “production app” in my opinion

Liveness → Is the app still running?

- Default is “process is running”
- Possible that the process can be running but not working correctly
- Good to define, might not be 100% necessary

*These can sometimes be the same endpoint, but not always*

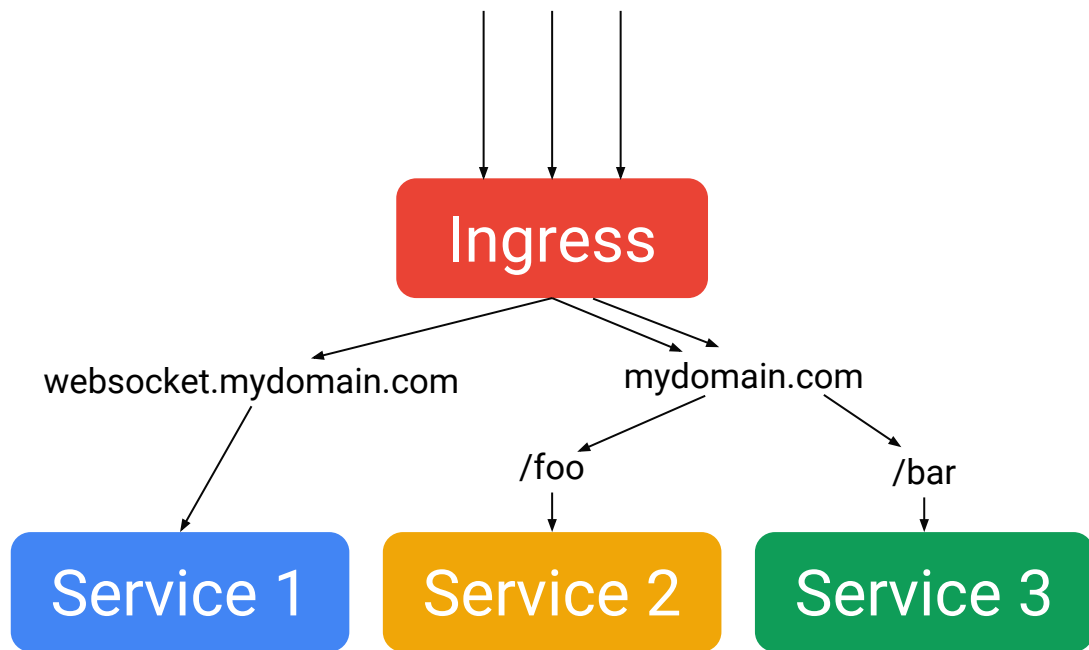


# Services

Don't always use `type: LoadBalancer`

Ingress is great





`type: NodePort` can be “good enough”

Use Static IPs. They are free\*!

```
$ gcloud compute addresses create ingress --global
...
$ gcloud compute addresses create myservice --region=us-west1
Created ...
address: QQQ.ZZZ.YYY.XXX
...
$
```

```
apiVersion: v1
kind: Service
metadata:
  name: myservice
spec:
  type: LoadBalancer
  loadBalancerIP: QQQ.ZZZ.YYY.XXX
  ports:
    - port: 80
      targetPort: 3000
      protocol: TCP
  selector:
    name: myapp
```

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: myingress
  annotations:
    kubernetes.io/ingress.global-static-ip-name: "ingress"
spec:
  backend:
    serviceName: myservice
    servicePort: 80
```

# Map external services to internal ones

# External Services

## Hosted Database

```
kind: Service
apiVersion: v1
metadata:
  name: mydatabase
  namespace: prod
spec:
  type: ExternalName
  externalName: my.database.example.com
  ports:
    - port: 12345
```

## Database outside cluster but inside network

<pre>kind: Service apiVersion: v1 metadata:   name: mydatabase spec:   ports:     - protocol: TCP       port: 80       targetPort: 12345</pre>	<pre>kind: Endpoints apiVersion: v1 metadata:   name: mydatabase subsets:   - addresses:       - ip: 10.128.0.2     ports:       - port: 12345</pre>
--	--



# Application Architecture

# Use Helm Charts



ALL downstream dependencies are unreliable

Make sure your microservices aren't too micro

# Use a “Service Mesh”



<https://github.com/istio/istio>



<https://github.com/linkerd/linkerd>

# Use a PaaS?





# Cluster Management

# Use Google Container Engine 🐳



# Resources, Anti-Affinity, and Scheduling

# Node Affinity

hostname

zone

region

instance-type

os

arch

custom!

# Node Taints / Tolerations

special hardware  
dedicated hosts  
etc

# Pod Affinity / Anti-Affinity

hostname

zone

region

# Use Namespaces to split up your cluster

# Role Based Access Control

# Unleash the Chaos Monkey

# More Resources

- <http://blog.kubernetes.io/2016/08/security-best-practices-kubernetes-deployment.html>
- <https://github.com/gravitational/workshop/blob/master/k8sprod.md>
- <https://nodesource.com/blog/8-protips-to-start-killing-it-when-dockerizing-node-js/>
- <https://www.ianlewis.org/en/using-kubernetes-health-checks>
- <https://www.linux.com/learn/rolling-updates-and-rollbacks-using-kubernetes-deployments>
- <https://kubernetes.io/docs/api-reference/v1.6/>





Questions?

What best practices do you have?